

Xiang (Emily) Zuo, PhD

220 Darryl Dr., Campbell, CA 95008, USA

+1 (408) 872-2860 | xiangzuo2012@gmail.com | <https://www.linkedin.com/in/emilyzuo/>

SUMMARY

AI/ML product leader at Teradata, building and scaling enterprise ML and generative AI products for **10+** years. Lead a **33-**person org that shipped Teradata's first GenAI product — an enterprise vector store now GA across 3 platforms. Combine deep LLM/multi-agent expertise with GTM and cross-functional leadership to ship revenue-relevant products.

EDUCATION

Stanford Ignite, Stanford University Graduate School of Business (GSB)	Jan. 2024 - Mar. 2024
Ph.D. in Computer Science, University of South Florida, Florida (USF)	2010 - 2016
M.S. in Computer Science, Huazhong University of Science and Technology (HUST), China	2006 - 2009
B.S. in Computer Science, Huazhong University of Science and Technology (HUST), China	2000 - 2004

WORK EXPERIENCE

Senior Technical Engineering Manager <i>Teradata</i> Santa Clara, CA	October 2023 - Now
- Lead a 33-person AI/ML team (6 PhDs), scaled from 15, across 6 products (up from 3) - Shipped Teradata's first GenAI product, an enterprise vector store — GA on 3 platforms, 12+ customers in active POC - Architect Tera, Teradata's multi-agent system, exposing all products via MCP servers and built-in agent skills — cutting token usage 5–30% and latency - Drove SLM strategy agents, retaining a high-end customer evaluating migration to Databricks via lower-cost inference - Set product roadmap and represent AI/ML in monthly product strategy reviews, aligning technical direction with GTM and customer priorities - Direct LLM-as-judge evaluation and production monitoring practices (LangSmith, DeepEval, MLflow) across the team's product line	
Staff AI Engineer <i>Teradata</i> Santa Clara, CA	March 2019 - October 2023
- Lead design and development of TeraAnalyticChat, a conversational AI product delivering NL Q&A and text-to-SQL generation over in-database analytic functions - Fine-tuned OpenLLaMA into a Teradata-branded LLM via PEFT (qLoRA, Adapter) and RAG over synthetic, public, and internal data, pairing a custom vector search engine for low-latency inference	
Senior Machine Learning Engineer <i>Teradata</i> Santa Clara, CA	Aug. 2016 - Feb. 2019
- Led end-to-end research, design, and development of 8 statistical and ML methods — XGBoost, TF-IDF, ANOVA, SMOTE, etc. — on Teradata's MPP distributed framework, expanding the in-database analytics function library - Serve as the POC site manager for enterprise customers including Verizon and AT&T	

SKILLS

Product & Leadership: Product Roadmap & Strategy, Org Building (15→33), Cross-functional Leadership, GTM Strategy, Competitive Positioning, Enterprise Stakeholder Management
Generative & Agentic AI: Agentic RAG, LLM Fine-Tuning (LoRA/QLoRA, PEFT), Prompt Engineering, Instruction Tuning, Multi-Agent Systems, Agent Skills & Harness Engineering, Model Context Protocol (MCP), Tool Use / Function Calling, CrewAI, Claude Agent SDK
ML/AI Frameworks: PyTorch, Hugging Face Transformers, LangChain, LangGraph, Scikit-learn
Evaluation & Observability: LM-as-Judge, LangSmith, DeepEval, MLFlow
Inference & Optimization: Quantization, Distillation, Pruning, vLLM, SGLang
Vector Databases & Retrieval: Pinecone, pgvector, Milvus, Hybrid Search, Reranking

HONORS AWARDS & PATENTS

Teradata 2024 Team Recognition Award, Company-wide awards	2024
Teradata 2024 Individual Recognition Award, Company-wide awards	2024
An Open-Source Alternative to ChatGPT, Teradata Hack Day 2nd Prize (2nd among 83)	June 2023
Hyperparameter Tuner for Vantage Analytics, Teradata Hack Day 1st Prize (Top 1 among 60)	June 2021
Patent: "Performing Hypersegmentation Analytics at Scale", Teradata Patent	Dec 2020